

DISTRIBUTED MESSAGE BROKER

ADVANCED NETWORK ARCHITECTURES AND
WIRELESS SYSTEMS

LEONARDO POGGIANI

CLARISSA POLIDORI

BRUNO AUGUSTO CASU PEREIRA DE SOUSA



SPECIFICATIONS




Figure 1 -

Computing nodes (either clients or servers) are connected through an **SDN-based network**.

The network acts as **a distributed message broker**, allowing computing nodes to *subscribe* or *publish* messages to resources.

Each *resource* is represented through a (virtual) IP address. Users can configure the system through a **REST-based Northbound interface**.

FUNCTIONALITIES

The system should:

- allow a user to **create** a resource;
- allow a user to **subscribe** a computing node to a resource;
- allow a computing node to **publish** a message on a resource;
- whenever a message is published on a resource, **deliver** it to every subscriber of the corresponding resource.

REST INTERFACE


http://controller_ip:8080/db/subscribers/RESOURCE_IP/json
http://controller_ip:8080/db/resources/json



REST INTERFACE

- Endpoint */resources*:
 - **GET**: Retrieves the list of registered resources.
 - **POST**: Creates a resource assigning to it a virtual address.
 - **DELETE**: Removes a resource from the registered resources.
- Endpoint */subscribers*:
 - **GET**: Retrieves the list of users subscribed to a resource.
 - **POST**: Adds a user to the list of subscribed users of a resource.
 - **DELETE**: Removes a user from the list of subscribed users to a resource.

DISTRIBUTED MESSAGE BROKER



DESIGN CHOICE

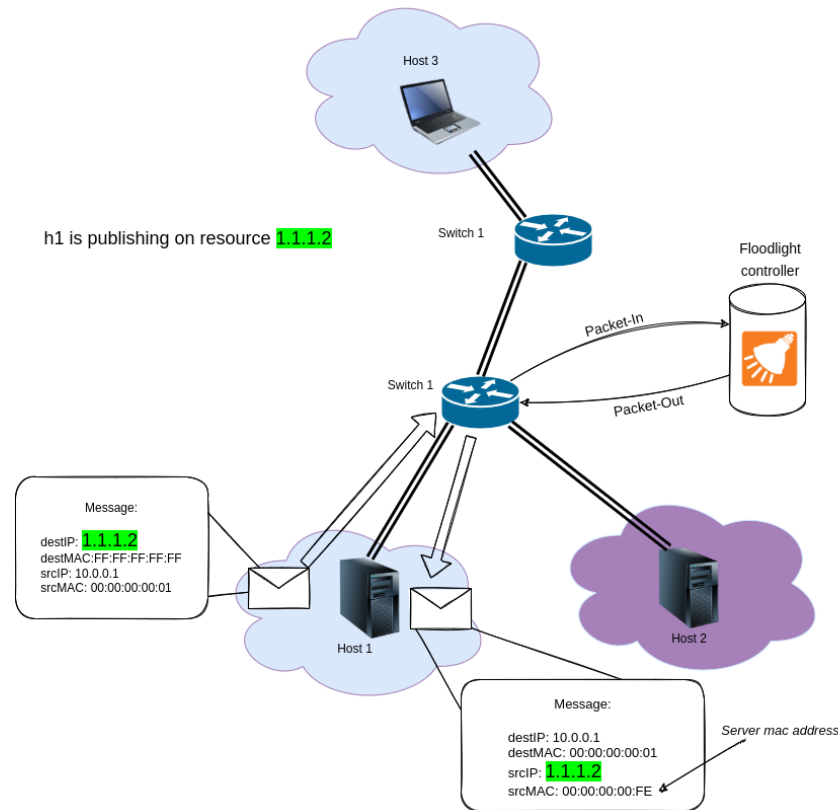
- **Only Packet-In and Packet-Out** instead of *flow-mods*
- **Only UDP traffic**
- ARP packets addressed to **virtual resources**
- **Custom forwarding module** instead of the *default one*.

UNWANTED BEHAVIOR

In our network we don't want to be possible that:

- A user send a message to another user using the **real (physical) destination address**.
- A user receive a message from a resource to which it is **not subscribed**.
- A user publish messages on a resource to **which he is subscribed**.

ARP REQUESTS



Example: h1 try to publish a message on resource 1.1.1.2

This originate a *packet-in* to the **Floodlight controller** who generates a *packet-out* containing an ARP REPLY message.

This ARP REPLY contains the source MAC address corresponding to the *SERVER_MAC address* and the source IP address equal to the resource address.

h1 is publishing on resource 1.1.1.2

Host 1

Host 2

Host 3

Host 4

Switch 1

Switch 2

Floodlight controller

Message:

destIP: 1.1.1.2
destMAC: 00:00:00:00:FE
srcIP: 10.0.0.1
srcMAC: 00:00:00:00:01

Message:

destIP: 10.0.0.2
destMAC: 00:00:00:00:02
srcIP: 1.1.1.2
srcMAC: 00:00:00:00:FE

Message:

destIP: 10.0.0.4
destMAC: 00:00:00:00:04
srcIP: 1.1.1.2
srcMAC: 00:00:00:00:FE

destIP not in resources IP, just forward

Packet-In

Packet-Out

This originate a *packet-in* to the **Floodlight controller** who generates a *packet-out* performing *two output actions*.

The output actions are redirected on the output port where they can reach (with the shortest path possible) their destinations.

TESTING



TESTING OBJECTIVES

- Tests will evaluate the system implemented features and performance.
- Two scenarios were developed, using the **mininet** software to provide a virtual network topology. For tracing the sent packets, **Wireshark** was used to capture the messages.
- From the virtual network, the available hosts will be used to run simple python scripts to test the nodes communication and message flow, using the developed Broker system.

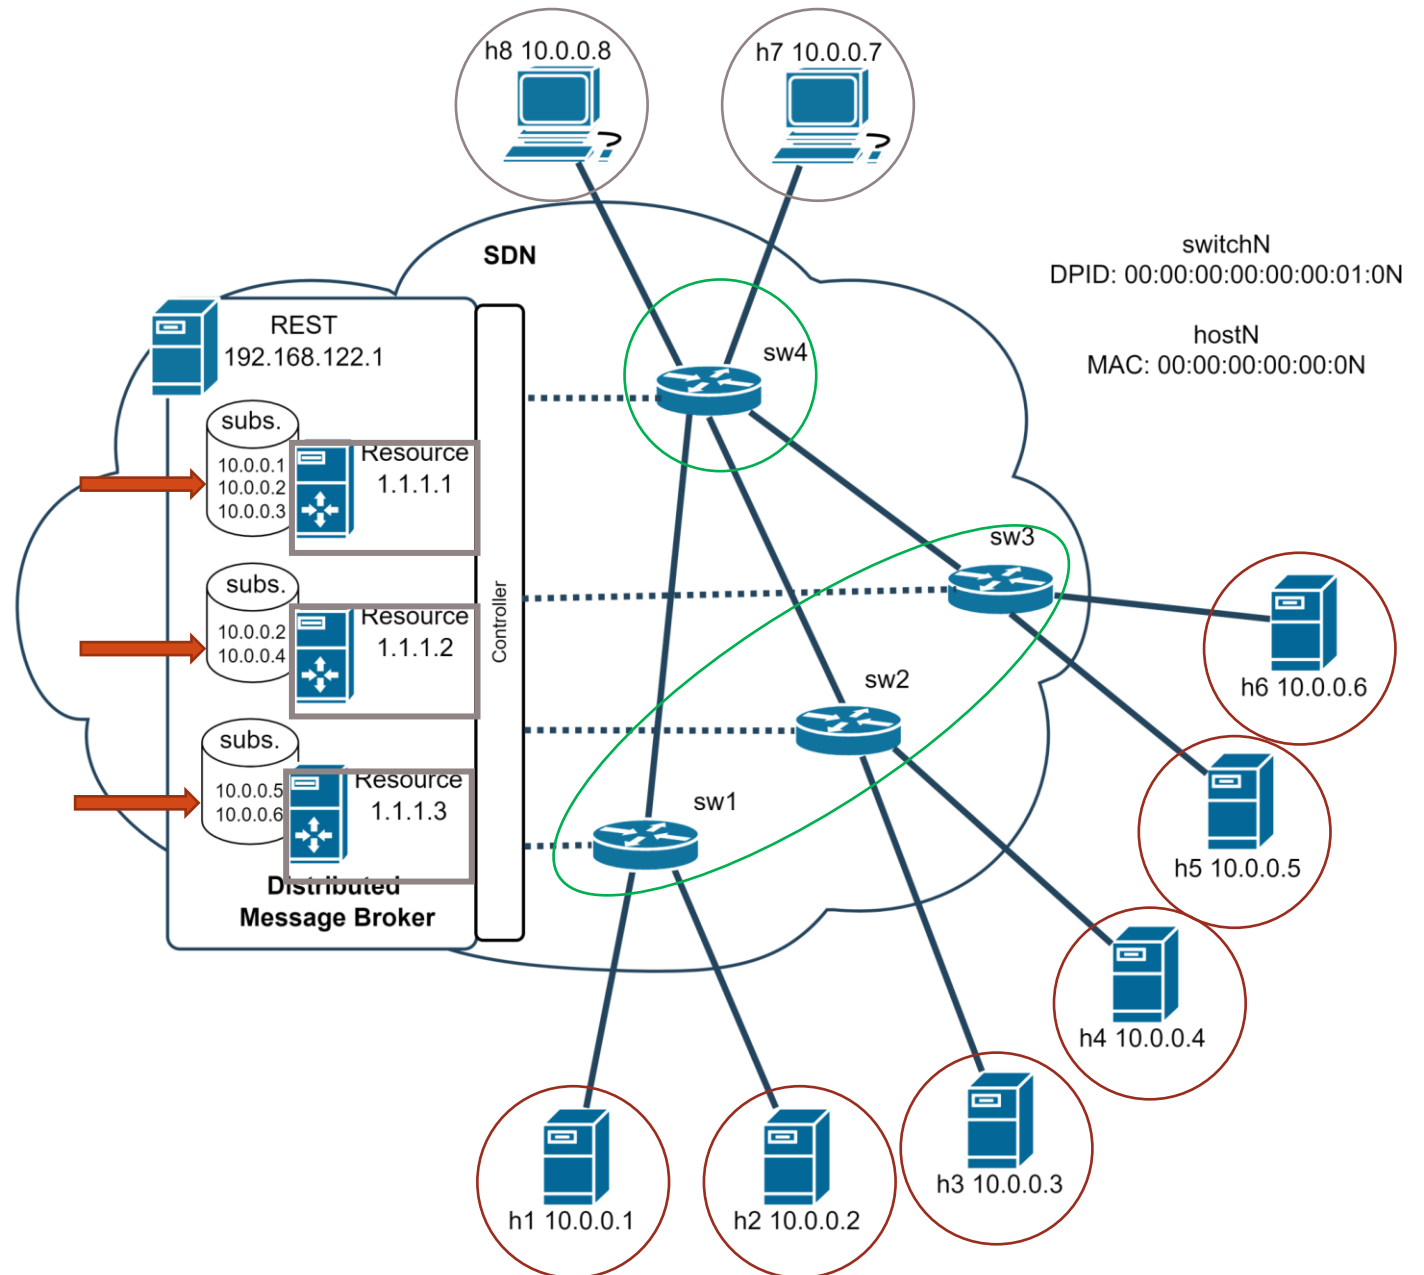
TEST TOPOLOGY

- Creating resources using the Broker REST interface

```
>>> usr.create_resource()
REST server response:
{"message":"Resource created, address: 1.1.1.1"}
>>> usr.create_resource()
REST server response:
{"message":"Resource created, address: 1.1.1.2"}
>>> usr.create_resource()
REST server response:
{"message":"Resource created, address: 1.1.1.3"}
```

- Subscribing the computing nodes using the Broker REST interface

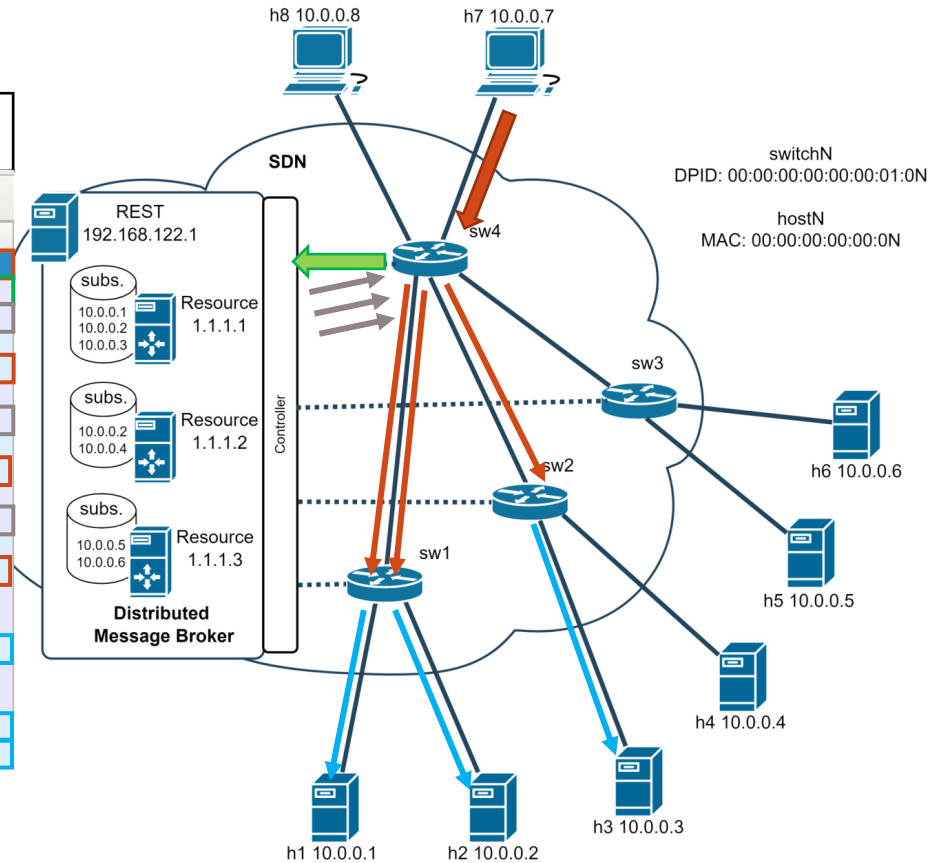
```
{ "message": "Resource created, address: 1.1.1.3" }
>>> usr.subscribe_cn("1.1.1.1", "10.0.0.1", "00:00:00:00:00:01")
REST server response:
{ "message": "Subscription successful" }
>>> usr.subscribe_cn("1.1.1.1", "10.0.0.2", "00:00:00:00:00:02")
REST server response:
{ "message": "Subscription successful" }
>>> usr.subscribe_cn("1.1.1.1", "10.0.0.3", "00:00:00:00:00:03")
REST server response:
{ "message": "Subscription successful" }
>>> usr.subscribe_cn("1.1.1.2", "10.0.0.2", "00:00:00:00:00:02")
REST server response:
{ "message": "Subscription successful" }
```




TESTING: PUBLISHING MESSAGES IN THE NETWORK

- Node h7 publishing on resource 1.1.1.1

No.	Time	Source	Destination	Protocol	Length	Info
2019...	72.939369990	10.0.0.7	1.1.1.1	UDP	57	59715 → 5005 Len=13
2019...	72.939796620	127.0.0.1	127.0.0.1	OpenFlow	165	Type: OFPT_PACKET_IN
2019...	72.944644857	127.0.0.1	127.0.0.1	OpenFlow	195	Type: OFPT_PACKET_OUT
2019...	72.944819838	10.0.0.7	10.0.0.3	UDP	57	59715 → 5005 Len=13
2019...	72.944822847	10.0.0.7	10.0.0.3	UDP	57	59715 → 5005 Len=13
2019...	72.945090065	127.0.0.1	127.0.0.1	OpenFlow	165	Type: OFPT_PACKET_IN
2019...	72.945545464	127.0.0.1	127.0.0.1	OpenFlow	195	Type: OFPT_PACKET_OUT
2019...	72.945673683	10.0.0.7	10.0.0.2	UDP	57	59715 → 5005 Len=13
2019...	72.945675517	10.0.0.7	10.0.0.2	UDP	57	59715 → 5005 Len=13
2019...	72.945859966	127.0.0.1	127.0.0.1	OpenFlow	165	Type: OFPT_PACKET_IN
2019...	72.946101595	127.0.0.1	127.0.0.1	OpenFlow	195	Type: OFPT_PACKET_OUT
2019...	72.946193626	10.0.0.7	10.0.0.1	UDP	57	59715 → 5005 Len=13
2019...	72.946195474	10.0.0.7	10.0.0.1	UDP	57	59715 → 5005 Len=13
2019...	72.946291817	127.0.0.1	127.0.0.1	OpenFlow	165	Type: OFPT_PACKET_IN
2019...	72.947363708	127.0.0.1	127.0.0.1	OpenFlow	163	Type: OFPT_PACKET_OUT
2019...	72.947615201	10.0.0.7	10.0.0.3	UDP	57	59715 → 5005 Len=13
2019...	72.951502856	127.0.0.1	127.0.0.1	OpenFlow	163	Type: OFPT_PACKET_OUT
2019...	72.951605209	127.0.0.1	127.0.0.1	OpenFlow	163	Type: OFPT_PACKET_OUT
2019...	72.951725004	10.0.0.7	10.0.0.2	UDP	57	59715 → 5005 Len=13
2019...	72.951776762	10.0.0.7	10.0.0.1	UDP	57	59715 → 5005 Len=13

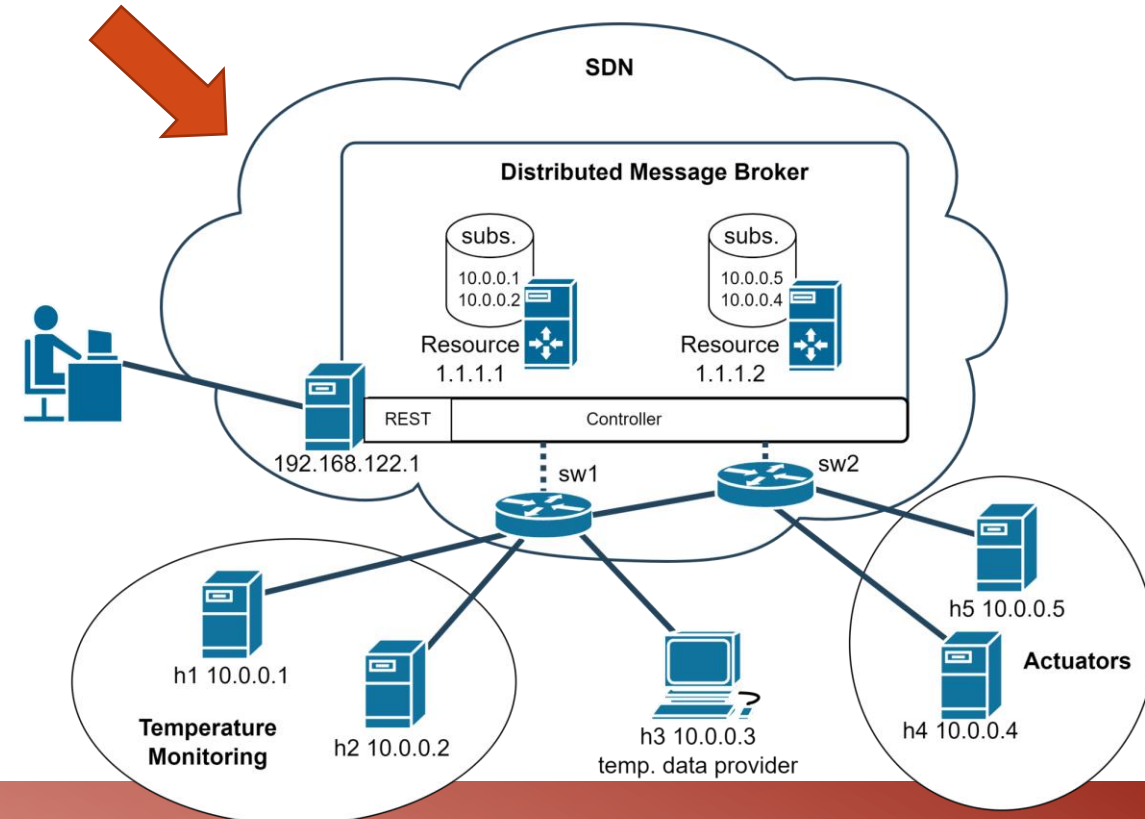


SENSOR APPLICATION IN LLN



Temperature Monitoring System

- A second **mininet** topology was implemented, and a simple python application was introduced in the available computing nodes, providing a simulated temperature monitoring system
- The system components are:
 - Data Provider node (acting as a publisher)
 - Monitoring Nodes (pool of servers subscribed to the monitoring topic)
 - Actuator Nodes (pool of server subscribed to the actuators topic)

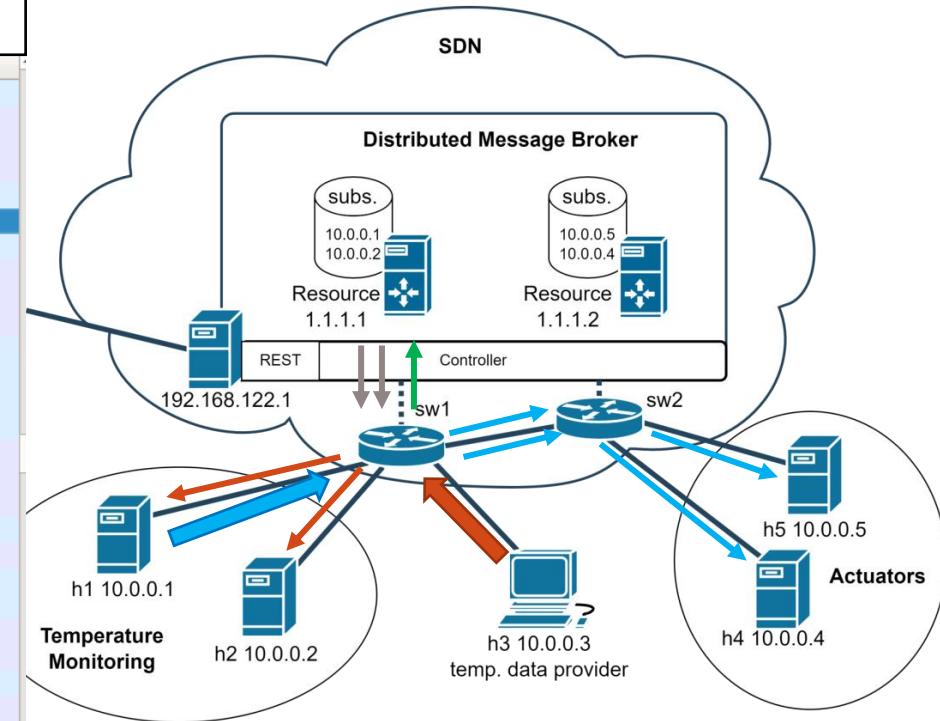


TESTING: SENSOR THRESHOLD CROSSED

- Node h3 (data source) publishing on resource 1.1.1.1 and h1 (monitoring node) publishing on resource 1.1.1.2

No.	Time	Source	Destination	Protocol	Length	Info
62606	22.050317585	10.0.0.3	1.1.1.1	UDP	49	54237 → 5005 Len=5
62607	22.050537994	127.0.0.1	127.0.0.1	OpenFlow	157	Type: OFPT_PACKET_IN
62608	22.051121494	127.0.0.1	127.0.0.1	OpenFlow	187	Type: OFPT_PACKET_OUT
62609	22.051165724	127.0.0.1	127.0.0.1	OpenFlow	187	Type: OFPT_PACKET_OUT
62612	22.051281218	10.0.0.3	10.0.0.2	UDP	49	54237 → 5005 Len=5
62614	22.051349774	10.0.0.3	10.0.0.1	UDP	49	54237 → 5005 Len=5
62619	22.052020031	10.0.0.1	1.1.1.2	UDP	51	5005 → 5005 Len=7
62620	22.052245571	127.0.0.1	127.0.0.1	OpenFlow	159	Type: OFPT_PACKET_IN
62621	22.052668691	127.0.0.1	127.0.0.1	OpenFlow	189	Type: OFPT_PACKET_OUT
62622	22.052750418	127.0.0.1	127.0.0.1	OpenFlow	189	Type: OFPT_PACKET_OUT
62623	22.052827336	10.0.0.1	10.0.0.5	UDP	51	5005 → 5005 Len=7
62624	22.052829542	10.0.0.1	10.0.0.5	UDP	51	5005 → 5005 Len=7
62628	22.053226558	10.0.0.1	10.0.0.4	UDP	51	5005 → 5005 Len=7
62629	22.053227797	10.0.0.1	10.0.0.4	UDP	51	5005 → 5005 Len=7
62630	22.053414340	127.0.0.1	127.0.0.1	OpenFlow	159	Type: OFPT_PACKET_IN
62631	22.053438840	127.0.0.1	127.0.0.1	OpenFlow	159	Type: OFPT_PACKET_IN
62633	22.054494817	127.0.0.1	127.0.0.1	OpenFlow	157	Type: OFPT_PACKET_OUT
62634	22.054635784	10.0.0.1	10.0.0.5	UDP	51	5005 → 5005 Len=7
62635	22.056287497	127.0.0.1	127.0.0.1	OpenFlow	157	Type: OFPT_PACKET_OUT
62637	22.056485450	10.0.0.1	10.0.0.4	UDP	51	5005 → 5005 Len=7
63524	22.552299270	10.0.0.3	1.1.1.1	UDP	49	51889 → 5005 Len=5
63525	22.553059925	127.0.0.1	127.0.0.1	OpenFlow	157	Type: OFPT_PACKET_IN
63532	22.555215930	127.0.0.1	127.0.0.1	OpenFlow	187	Type: OFPT_PACKET_OUT
63533	22.555499874	10.0.0.3	10.0.0.2	UDP	49	51889 → 5005 Len=5
63534	22.558086832	127.0.0.1	127.0.0.1	OpenFlow	187	Type: OFPT_PACKET_OUT
63536	22.558303029	10.0.0.3	10.0.0.1	UDP	49	51889 → 5005 Len=5
64977	23.053372909	10.0.0.3	1.1.1.1	UDP	49	40976 → 5005 Len=5
64978	23.053554145	127.0.0.1	127.0.0.1	OpenFlow	157	Type: OFPT_PACKET_IN


▶ Frame 62614: 49 bytes on wire (392 bits), 49 bytes captured (392 bits) on interface 0
▶ Linux cooked capture
▶ Internet Protocol Version 4, Src: 10.0.0.3, Dst: 10.0.0.1
▶ User Datagram Protocol, Src Port: 54237, Dst Port: 5005
▼ Data (5 bytes)
Data: 32342e3733
[Length: 5]



```
Temperature Sensor Reading (C): 24.88
Temperature Sensor Reading (C): 24.85
Temperature Sensor Reading (C): 24.12
Temperature Sensor Reading (C): 25.15
Temperature Sensor Reading (C): 24.73
Moving Average (C): 24.75
```


Threshold Crossed - Trigger Actuators

CONCLUSIONS AND PERFORMANCE EVALUATION




• System scalability

- IoT systems can be designed with multiple clusters of sensor nodes, that transmits its local readings to the monitoring centers.
- The Broker system then provides a better way to handle an increased volume of messages that needs to reach the aggregator node, reducing the traffic in the Radio Link.
- The plot shows the expected number of packets transmitted in the Radio Link when using the Broker system and when using P2P communication, with different system workloads.



Estimation of packet rate in the radio link by the number of deployed sensor data generators




Rate **R** is the amount of pkts generated by the Data Providers (500 ms temp. Reading interval results in $R=2\text{pkts/s}$)

THANKS FOR THE ATTENTION!

- Project github repo: <https://github.com/leonardopoggiani/ANAWS-project>
- IoT images extracted from Prof. Anastasi lecturing material

BACKUP SLIDES



DETAILED TESTING OBJECTIVES

- All the subscribers to a resource receive the messages published on it
- Non-subscribers do not see message published on the resource
- A host can subscribe to multiple resources
- A host that unsubscribes no longer receives messages published to that resource
- When a resource is deleted, no host is subscriber anymore and receives nothing
- A host can not send a message to another host directly
- A host cannot publish to a resource to which it is subscribed

TESTING: PUBLISHING MESSAGES IN THE NETWORK

- Node h1 publishing on resource 1.1.1.2

Capturing from h4-eth0

No.	Time	Source	Destination	Protocol	Length	Info
1	0.00000000	10.0.0.1	10.0.0.4	UDP	55	41290 → 5005 Len=13
2	3.599517237	36:f1:ba:35:a5:94		LLDP Multicast	LLDP	75 TTL = 120
3	3.628440663	36:f1:ba:35:a5:94		Broadcast	0x8942	83 Ethernet II
4	18.738646796	36:f1:ba:35:a5:94		LLDP Multicast	LLDP	75 TTL = 120
5	18.777375090	36:f1:ba:35:a5:94		Broadcast	0x8942	83 Ethernet II

Capturing from h2-eth0

No.	Time	Source	Destination	Protocol	Length	Info
1	0.00000000	10.0.0.1	10.0.0.2	UDP	55	41290 → 5005 Len=13
2	3.597364899	ea:3b:3d:41:01:c0		LLDP Multicast	LLDP	75 TTL = 120
3	3.626080158	ea:3b:3d:41:01:c0		Broadcast	0x8942	83 Ethernet II
4	18.73555450	ea:3b:3d:41:01:c0		LLDP Multicast	LLDP	75 TTL = 120
5	18.775013834	ea:3b:3d:41:01:c0		Broadcast	0x8942	83 Ethernet II

Capturing from h1-eth0

No.	Time	Source	Destination	Protocol	Length	Info
1	0.00000000	00:00:00 00:00:01	Broadcast	ARP	42	Who has 1.1.1.2? Tell 10.0.0.1
2	0.000835686	00:00:00 00:00:fe	00:00:00 00:00:01	ARP	42	1.1.1.2 is at 00:00:00:00:00:fe
3	0.000841811	10.0.0.1	1.1.1.2	UDP	55	41290 → 5005 Len=13
4	3.606403899	a2:f5:46:c0:09:96		LLDP Multicast	LLDP	75 TTL = 120
5	3.634262993	a2:f5:46:c0:09:96		Broadcast	0x8942	83 Ethernet II
6	18.747517475	a2:f5:46:c0:09:96		LLDP Multicast	LLDP	75 TTL = 120
7	18.783221227	a2:f5:46:c0:09:96		Broadcast	0x8942	83 Ethernet II

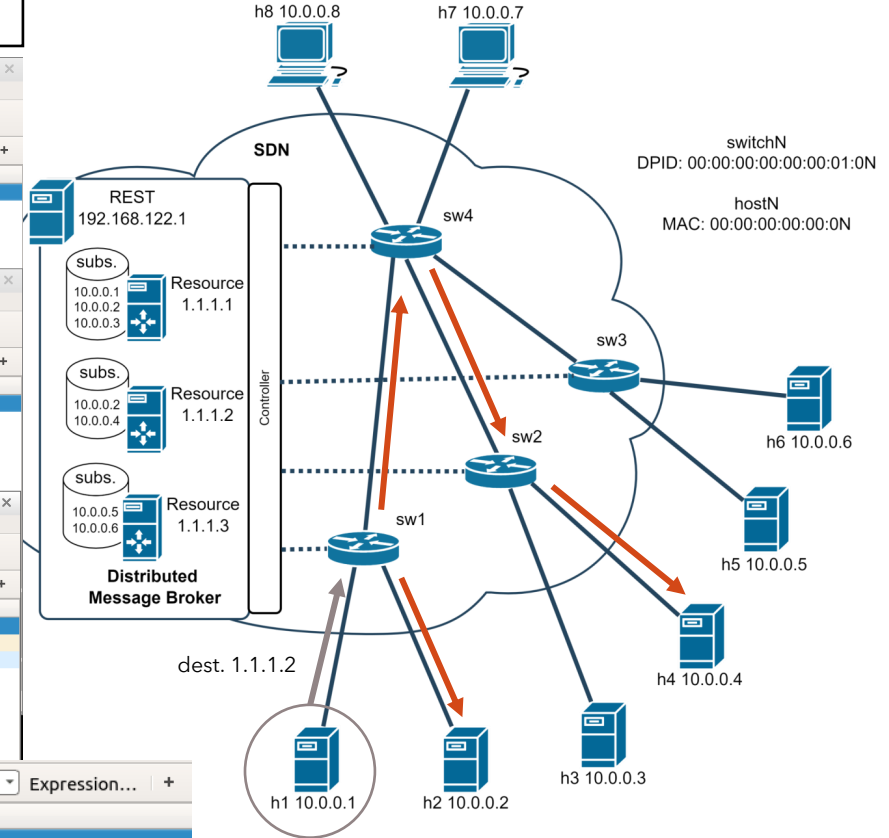
Received UDP msg
(h4 and h2 interfaces
- subscribers)

ARP resolution
provided by the Broker
module

Sent UDP msg (h1
interface - publisher)

All interfaces

No.	Time	Source	Destination	Protocol	Length	Info
5463...	2058.4314177...	10.0.0.1	1.1.1.2	UDP	57	42031 → 5005 Len=13
5463...	2058.4316092...	127.0.0.1	127.0.0.1	OpenFlow	165	Type: OFPT_PACKET_IN
5464...	2058.9110114...	127.0.0.1	127.0.0.1	OpenFlow	195	Type: OFPT_PACKET_OUT
5464...	2058.9111133...	127.0.0.1	127.0.0.1	OpenFlow	195	Type: OFPT_PACKET_OUT
5464...	2058.9111894...	10.0.0.1	10.0.0.4	UDP	57	42031 → 5005 Len=13
5464...	2058.9111913...	10.0.0.1	10.0.0.4	UDP	57	42031 → 5005 Len=13
5464...	2058.9113082...	10.0.0.1	10.0.0.2	UDP	57	42031 → 5005 Len=13
5464...	2058.9114392...	127.0.0.1	127.0.0.1	OpenFlow	165	Type: OFPT_PACKET_IN
5464...	2058.9134508...	127.0.0.1	127.0.0.1	OpenFlow	163	Type: OFPT_PACKET_OUT
5464...	2058.9135862...	10.0.0.1	10.0.0.4	UDP	57	42031 → 5005 Len=13
5464...	2058.9135879...	10.0.0.1	10.0.0.4	UDP	57	42031 → 5005 Len=13
5464...	2058.9137133...	127.0.0.1	127.0.0.1	OpenFlow	165	Type: OFPT_PACKET_IN
5465...	2059.0919025...	127.0.0.1	127.0.0.1	OpenFlow	163	Type: OFPT_PACKET_OUT
5465...	2059.0920154...	10.0.0.1	10.0.0.4	UDP	57	42031 → 5005 Len=13



CONFIG OTHER TESTS

